

# Text Classification for Economics

A Practical Introduction to Python & Machine Learning

---

HKU ECON6087 Tutorial

2026 年 2 月 5 日

# Objective

- **Goal:** Classify news articles into categories (World, Sports, Business, Sci/Tech).
- **Tools:** Python and its data science libraries.
- **Method:** K-Nearest Neighbors (KNN).
- **Key Learning:** How specific text representations (Bag of Words vs. TF-IDF) affect model performance.

## Python Libraries

---

# The Python Data Stack

We use three main libraries—think of them as your digital toolbox:

**datasets (Hugging Face)**: Like a library catalog. It lets us easily download and access standard datasets without manual searching.

**pandas (Data Manipulation)**: The "Excel of Python". It handles data tables (DataFrames), allowing us to load, view, and clean data efficiently.

**scikit-learn (Machine Learning)**: The standard toolkit for machine learning. It contains the algorithms (like KNN) and tools to transform text into numbers.

## The Data

---

# The Dataset: AG News

We use the **AG News** dataset used in the tutorial.

- **Content:** News headlines and short descriptions.
- **Classes:** 4 Topics
  - 1. World
  - 2. Sports
  - 3. Business
  - 4. Sci/Tech
- **Size:**
  - Training Set: 120,000 articles (used to teach the model).
  - Testing Set: 7,600 articles (used to evaluate performance).

## Step 1: Loading Data

---

## Step 1: Loading and Saving Data

We first load the data from the cloud and save it locally as CSV files.

```
1 from datasets import load_dataset
2
3 # 1. Download data
4 dataset = load_dataset("ag_news")
5
6 # 2. Save as CSV (for Pandas to read later)
7 for split in dataset.keys():
8     # dataset[split] is the data (train or test)
9     # .to_csv() saves it to a file
10    dataset[split].to_csv(f"{split}.csv")
```

*Result: We now have ‘train.csv’ and ‘test.csv’.*

## Concept: K-Nearest Neighbors

---

# Concept: K-Nearest Neighbors (KNN)

## Intuition

"Tell me who your neighbors are, and I'll tell you who you are."

To classify a new document:

1. We place the new document in the "space" of all known documents.
2. We find the **K** closest documents (neighbors).
3. We look at their labels (topics).
4. We assign the majority label to the new document.

**Challenge for Text:** How do we measure distance between text?  
We must turn words into numbers.

## Step 2: Bag of Words

---

## Text Representation 1: Bag of Words (BoW)

The simplest way to turn text into numbers.

- We list all unique words in the dataset (the vocabulary).
- For each document, we simply **count** how many times each word appears.

### 例

"The economy grows" → {'the': 1, 'economy': 1, 'grows': 1, 'football': 0}

**Pros:** Simple.

**Cons:** Ignores grammar. Common words like "the" appear frequently but have little meaning.

# Implementing Bag of Words

In Python, CountVectorizer creates the Bag of Words matrix.

```
1 from sklearn.feature_extraction.text import CountVectorizer  
2  
3 # Create a "counter" that looks at the top 5000 words  
4 # stop_words='english' removes common words like "the", "is", "  
# at"  
5 vectorizer = CountVectorizer(stop_words='english', max_features  
# =5000)  
6  
7 # Learn the vocabulary and count words in training data  
8 X_train_bow = vectorizer.fit_transform(train_df['text'])  
9 # Count words in test data (using same vocabulary)  
10 X_test_bow = vectorizer.transform(test_df['text'])
```

## Results: KNN with Bag of Words

```
1 from sklearn.neighbors import KNeighborsClassifier  
2  
3 # Create the model (look at 5 nearest neighbors)  
4 knn = KNeighborsClassifier(n_neighbors=5)  
5  
6 # Train: Model memorizes the training data  
7 knn.fit(X_train_bow, train_df['label'])  
8  
9 # Predict and Evaluate  
10 accuracy = knn.score(X_test_bow, test_df['label'])  
11 print(f"Accuracy: {accuracy}")
```

### Result

Accuracy  $\approx$  72.5%

## Step 3: TF-IDF

---

## Text Representation 2: TF-IDF

Term Frequency - Inverse Document Frequency.

A smarter way to count.

- **TF (Term Frequency)**: How often word  $w$  appears in document  $d$ .
- **IDF (Inverse Document Frequency)**: How rare is word  $w$  across *all* documents?

**Idea:** If a word appears in *every* document (e.g., "said", "today"), it's not useful for classification. We lower its weight. If a word is rare (e.g., "touchdown", "deficit"), it gets a high weight.

# Implementing TF-IDF

Changing the code is easy: swap `CountVectorizer` for `TfidfVectorizer`.

```
1 from sklearn.feature_extraction.text import TfidfVectorizer  
2  
3 # Use TF-IDF instead of simple counts  
4 tfidf_vectorizer = TfidfVectorizer(stop_words='english',  
5                                   max_features=5000)  
6  
7 # Transform data  
8 X_train_tfidf = tfidf_vectorizer.fit_transform(train_df['text'])  
X_test_tfidf = tfidf_vectorizer.transform(test_df['text'])
```

## Results: KNN with TF-IDF

We train the exact same KNN model, just using the new weighted data.

```
1 # Train KNN on TF-IDF data
2 knn_tfidf = KNeighborsClassifier(n_neighbors=5)
3 knn_tfidf.fit(X_train_tfidf, train_df['label'])
4
5 # Evaluate
6 accuracy_tfidf = knn_tfidf.score(X_test_tfidf, test_df['label'])
7 print(f"Accuracy: {accuracy_tfidf}")
```

### Result

Accuracy  $\approx$  89.0%

**Takeaway:** Weighting words by importance (TF-IDF) massively improves our simple neighbor-based classifier.

## Extension: Chinese Text

---

# Working with Chinese Data

**Challenge:** English words are separated by spaces. Chinese words are not.

- English: "I love economics" → ["I", "love", "economics"]
- Chinese: "我爱经济学" → ?

**Solution:** Word Segmentation (Tokenization).

- We use a library called **jieba** ("Stutter" in Chinese) to cut sentences into words.
- "我爱经济学"  $\xrightarrow{jieba}$  ["我", "爱", "经济学"]

# Chinese Segmentation with Jieba

```
1 import jieba  
2  
3 # Example Function  
4 def segment_text(text):  
5     # jieba.cut returns a generator, we join it with spaces  
6     return " ".join(jieba.cut(text))  
7  
8 # Apply to a sentence  
9 print(segment_text("我爱经济学"))  
10 # Output: "我 爱 经济学"
```

Once the text is space-separated, we can use **CountVectorizer** and **TfidfVectorizer** exactly as before!

## Conclusion

---

## Summary

1. **Pandas** allows economists to handle large datasets easily.
2. **Scikit-Learn** provides plug-and-play machine learning tools.
3. Text must be converted to numbers before analysis.
4. **TF-IDF** is often superior to simple counting because it filters out "noise" and highlights unique keywords.
5. Even a simple algorithm like **KNN** can achieve high accuracy (~89%) with the right data representation.